

Why Scaling Large Language Models Alone May Not Lead to Artificial General Intelligence

Abdul Wahab
Independent Researcher

May 2026

Abstract

Artificial General Intelligence (AGI) remains one of the most ambitious goals in artificial intelligence research. Recent advances in transformer-based large language models have demonstrated powerful emergent capabilities, leading to debates regarding whether scaling alone may eventually produce AGI. However, despite impressive performance across reasoning and language tasks, current systems still exhibit major limitations in long-term planning, memory persistence, grounded understanding, and autonomous reasoning. This paper analyzes the architectural limitations of modern large language models and explores additional components that may be necessary for achieving AGI.

1 Introduction

Artificial General Intelligence (AGI) refers to systems capable of performing a broad range of cognitive tasks at or beyond human-level intelligence. Unlike narrow artificial intelligence systems, which are designed for specific domains, AGI aims to generalize knowledge and reasoning across diverse environments and tasks [4].

The rapid progress of transformer-based large language models has significantly accelerated discussions surrounding AGI. Transformer architectures introduced substantial improvements in language modeling and representation learning capabilities [10]. Recent frontier large language models such as GPT-4, Gemini, and Claude have demonstrated increasingly sophisticated capabilities in reasoning, coding, language understanding, and multimodal interaction [2, 3].

Despite these advances, major questions remain regarding whether scaling current architectures alone is sufficient for achieving true AGI. Current systems continue to struggle with persistent memory, long-term planning, grounded reasoning, and robust world understanding [1].

This paper investigates the limitations of scaling-based approaches and discusses possible architectural directions that may contribute toward the development of AGI systems.

2 Background

The development of artificial intelligence has progressed through multiple technological and conceptual stages over the past several decades. Early artificial intelligence systems were primarily rule-based expert systems designed to solve narrowly defined problems through manually engineered logic. Although these systems demonstrated success in structured environments, they lacked adaptability and were unable to generalize knowledge beyond predefined domains [9].

The emergence of machine learning significantly transformed the field by enabling systems to learn patterns directly from data rather than relying entirely on handcrafted rules. In particular, deep learning architectures based on artificial neural networks achieved major breakthroughs

in image recognition, speech processing, and natural language understanding [6]. Increasing computational power, large-scale datasets, and advances in parallel hardware accelerated the capabilities of these systems.

A major turning point in modern artificial intelligence research occurred with the introduction of transformer architectures. The transformer model, introduced in the paper “Attention Is All You Need,” replaced recurrent processing with self-attention mechanisms capable of modeling long-range dependencies efficiently [10]. This architectural innovation enabled substantial scaling in model size and training data, forming the foundation of modern large language models.

Large language models (LLMs) are trained on massive corpora of text data using self-supervised learning objectives, most commonly next-token prediction. During training, these models learn statistical representations of language, syntax, reasoning patterns, and contextual relationships. Scaling laws observed in recent years suggest that increasing model parameters, dataset size, and computational resources can produce emergent capabilities not explicitly programmed during training [5].

Recent frontier systems such as GPT-4, Gemini, Claude, and other multimodal architectures have demonstrated increasingly advanced performance across a broad range of tasks, including code generation, mathematical reasoning, summarization, translation, scientific assistance, and conversational interaction [2, 3]. Reinforcement Learning from Human Feedback (RLHF) has further improved alignment between model outputs and human expectations, contributing to more coherent and contextually relevant responses [7].

Despite these advances, the distinction between narrow artificial intelligence and Artificial General Intelligence remains significant. Narrow AI systems are optimized for specific tasks and operate primarily through statistical pattern recognition within learned distributions. AGI, by contrast, is generally understood as a system capable of robust reasoning, autonomous planning, transfer learning across domains, persistent memory, adaptability to unfamiliar environments, and generalized cognitive competence comparable to or exceeding human intelligence [4].

The rapid progress of large language models has generated substantial debate regarding whether scaling transformer-based architectures alone may eventually lead to AGI. Some researchers argue that sufficiently large models may continue to exhibit increasingly sophisticated emergent behaviors as computational scale increases. Others argue that current architectures fundamentally lack important components required for true general intelligence, including grounded world models, long-term memory systems, causal reasoning, embodiment, and autonomous objective formation [1].




Figure 1: Conceptual evolution of artificial intelligence systems from rule-based architectures toward potential Artificial General Intelligence.

This ongoing debate forms the foundation of contemporary AGI research and motivates the investigation presented in this paper.

3 Limitations of Current Large Language Models

Large language models have demonstrated remarkable progress across language understanding, reasoning, coding, and multimodal interaction tasks. However, despite these advances, current transformer-based systems still exhibit several fundamental limitations that challenge the possibility of achieving true Artificial General Intelligence through scaling alone [3, 1].

One of the most significant limitations of current LLMs is the absence of persistent long-term memory. Although models can process large context windows during inference, they do not possess stable memory structures comparable to human cognition. Information provided during interactions is typically constrained by context length limitations and session-based processing. As a result, models struggle to maintain consistent long-term understanding, evolving objectives, and persistent world representations across extended interactions [8].

Another major limitation involves grounded understanding. Current LLMs primarily learn statistical relationships between tokens from large-scale textual datasets rather than developing direct experiential models of the physical world. Consequently, model outputs may appear coherent while lacking genuine semantic understanding or causal reasoning. This often contributes to hallucinations, where models generate factually incorrect or fabricated information with high confidence [1].

Reasoning capabilities in modern language models also remain limited in important ways. Although chain-of-thought prompting and reinforcement learning techniques have improved performance on reasoning benchmarks, many systems continue to rely heavily on probabilistic pattern matching rather than robust symbolic or causal reasoning. Performance degradation frequently occurs when models encounter unfamiliar tasks, adversarial prompts, or highly compositional reasoning problems requiring multi-step planning [11].

Autonomous planning and objective formation present additional challenges. Current language models generally operate as reactive systems responding to user prompts rather than independently forming long-term goals or adaptive strategies. While recent agentic architectures attempt to extend model autonomy through tool usage, memory modules, and external planning systems, these capabilities often remain fragile and dependent on carefully engineered prompting frameworks [12].

Another important issue concerns computational efficiency and scalability. Modern frontier models require extremely large computational resources, massive datasets, and extensive energy consumption during both training and inference. Continued scaling may eventually encounter practical economic, environmental, and hardware limitations. Furthermore, increasing parameter count alone does not necessarily guarantee proportional improvements in reasoning robustness, factual reliability, or general intelligence [5].

Current LLMs also face challenges in interpretability and transparency. The internal decision-making processes of large neural networks remain difficult to fully understand, limiting the ability of researchers to diagnose failures, verify reasoning pathways, or ensure robust alignment with human intentions. This lack of interpretability becomes increasingly concerning as models are deployed in high-impact domains such as healthcare, scientific research, cybersecurity, and autonomous systems [1].

In addition to technical limitations, alignment and safety concerns represent major obstacles toward AGI development. Advanced systems may exhibit deceptive behavior, reward hacking, or unintended optimization strategies under certain training conditions. Ensuring reliable alignment between increasingly capable AI systems and human values remains an open research problem within the broader AGI community [7].

These limitations suggest that although scaling transformer architectures has produced substantial progress in artificial intelligence capabilities, current large language models may still lack several core components required for robust general intelligence.




Figure 2: Conceptual overview of major limitations in current large language model architectures.

4 Possible Paths Toward AGI

Although current large language models demonstrate impressive capabilities, many researchers argue that achieving Artificial General Intelligence will likely require architectures extending beyond large-scale next-token prediction systems [1, 4]. Several research directions have emerged as possible paths toward more general and adaptive intelligence, including memory-augmented architectures, symbolic reasoning systems, agentic AI frameworks, grounded learning approaches, and hybrid computational models [12, 8].

4.1 Better Memory Systems

One major limitation of current LLMs is the absence of persistent long-term memory [8]. Future AGI systems may require advanced memory architectures capable of storing, updating, and retrieving information across extended periods of interaction. Persistent memory could enable improved contextual continuity, adaptive learning, and long-term reasoning capabilities beyond the constraints of fixed context windows [1].

4.2 Improved Reasoning Mechanisms

Although modern language models exhibit increasingly sophisticated reasoning behaviors, many systems still struggle with causal inference, logical consistency, and complex multi-step planning [11, 1]. Future AGI architectures may integrate symbolic reasoning systems alongside neural networks to improve interpretability, logical robustness, and general problem-solving performance across unfamiliar domains [4].

4.3 Agentic AI Systems

Recent advances in agentic artificial intelligence have explored systems capable of autonomous planning, tool usage, self-reflection, and iterative decision-making [12, 8]. Rather than functioning solely as reactive conversational systems, future AGI architectures may involve autonomous agents capable of pursuing long-term objectives, adapting strategies, and interacting dynamically with external environments [1].

4.4 Grounded and Embodied Learning

Current LLMs primarily learn from textual datasets and therefore lack direct experiential understanding of the physical world [1]. Some researchers argue that future AGI systems may require grounded learning through multimodal interaction, environmental feedback, robotics, or simulated world models. Such approaches may contribute toward stronger causal understanding and more robust generalization capabilities [4].

4.5 Hybrid Architectures

A growing body of research suggests that AGI may ultimately require hybrid architectures combining multiple computational paradigms [4, 1]. Rather than relying exclusively on transformer-based language prediction, future systems may integrate memory modules, symbolic reasoning frameworks, reinforcement learning, planning systems, and multimodal world models into unified architectures capable of more generalized intelligence.

These research directions suggest that although scaling large language models has produced substantial advances in artificial intelligence, achieving true AGI will likely require broader architectural innovations extending beyond current transformer-based systems [5, 3].

5 Conclusion

The rapid advancement of transformer-based large language models has significantly accelerated progress in artificial intelligence research. Modern systems have demonstrated increasingly sophisticated capabilities across reasoning, language understanding, coding, and multimodal interaction tasks, leading to renewed debate surrounding the possibility of Artificial General Intelligence emerging through scaling alone.

However, despite these advances, current large language models continue to exhibit important limitations in long-term memory, grounded understanding, causal reasoning, autonomous planning, and interpretability. These limitations suggest that statistical language prediction alone may be insufficient for achieving robust general intelligence comparable to human cognition.

This paper has argued that future progress toward AGI will likely require broader architectural innovations extending beyond current transformer-based systems. Research directions involving persistent memory systems, symbolic reasoning, agentic architectures, grounded learning, and hybrid computational frameworks may contribute toward more adaptive and generalized intelligence.

Although the path toward AGI remains uncertain, the continued development of increasingly capable artificial intelligence systems will likely remain one of the most important scientific and technological challenges of the twenty-first century. Understanding both the capabilities and limitations of current AI architectures is therefore essential for guiding future research toward safe, interpretable, and genuinely general intelligent systems.

References

- [1] Yoshua Bengio. Towards conscious ai systems. *arXiv preprint arXiv:2308.08708*, 2023.
- [2] Tom B. Brown et al. Language models are few-shot learners. *Advances in Neural Information Processing Systems*, 2020.
- [3] Sebastien Bubeck et al. Sparks of artificial general intelligence: Early experiments with gpt-4. *arXiv preprint arXiv:2303.12712*, 2023.
- [4] Ben Goertzel. Artificial general intelligence: Concept, state of the art, and future prospects. *Journal of Artificial General Intelligence*, 2014.
- [5] Jared Kaplan et al. Scaling laws for neural language models. *arXiv preprint arXiv:2001.08361*, 2020.
- [6] Yann LeCun, Yoshua Bengio, and Geoffrey Hinton. Deep learning. *Nature*, 521(7553):436–444, 2015.
- [7] Long Ouyang et al. Training language models to follow instructions with human feedback. *Advances in Neural Information Processing Systems*, 2022.
- [8] Joon Sung Park et al. Generative agents: Interactive simulacra of human behavior. *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology*, 2023.
- [9] Stuart Russell and Peter Norvig. *Artificial Intelligence: A Modern Approach*. Prentice Hall, 2010.
- [10] Ashish Vaswani et al. Attention is all you need. *Advances in Neural Information Processing Systems*, 2017.
- [11] Jason Wei et al. Chain-of-thought prompting elicits reasoning in large language models. *arXiv preprint arXiv:2201.11903*, 2022.
- [12] Shunyu Yao et al. React: Synergizing reasoning and acting in language models. *arXiv preprint arXiv:2210.03629*, 2022.